

# main

December 7, 2025

```
[89]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import time
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import learning_curve
df_train = pd.read_csv('train.csv',index_col='id')
df_test = pd.read_csv('test.csv',index_col='id')
```

## 1 EXPLORING THE DATA

```
[90]: df_train.head()
```

```
[90]:      f1        f2        f3        f4        f5        f6        f7 \
id
1   0.813120 -0.317922  0.027886  2.966993  1.792626  1.639461  1.038417
2   1.187370  0.574752  0.480094  1.003127  1.743455  1.653923  0.289389
3   0.713207 -0.061996  0.423746  0.111901  1.763365  1.651579  1.026891
4   0.779316  0.687488  0.388627  1.343889  1.743500  1.641686 -0.670589
5   0.674119 -0.286842  0.386524  2.416830  1.787492  1.636464  1.782221


      f8        f9        f10       ...       f12        f13       f14        f15 \
id
1   -1.265553  2.737895  3.723351  ...  2.074784 -0.549459  0.588724  0.670610
2    1.758154  1.447072  0.014709  ...  0.904899  0.936627  0.692230  0.631463
3    1.283781  1.116797  0.291573  ...  1.581493  1.151487  0.213017  0.766878
4    0.135970  1.553194  2.046264  ...  1.995739  0.650913  0.925546  0.657044
5    0.722959  2.030776  0.550102  ...  0.485689 -1.000375  0.970572  0.900548

      f16       f17       f18       f19       f20     target
id
1   0.102525  1.963463 -0.349565 -0.294764 -0.206112 -18.759727
2   1.022266  2.211432 -0.318560  0.696366  0.693999  8.082223
3   1.003003  1.824628  1.248738      NaN  0.600792  2.961315
4   0.558577  1.915205  2.763816  1.665433  0.397841 -0.710178
5   0.638401  1.893995 -2.186096  1.838846  1.070571  5.328069
```

```
[5 rows x 21 columns]
```

```
[91]: df_train['target'].plot()
```

```
[91]: <Axes: xlabel='id'>
```



Notes:

Need something that can deal with outliers as the above graph shows

Additionally we will need an imputer to deal with outliers

```
[92]: df_train.isna().sum() #!only minor 74 total but only 400 examples so cant just  
→ get rid of them, imputer?
```

```
[92]: f1      5  
f2      2  
f3      3  
f4      3  
f5      7  
f6      2  
f7      4
```

```

f8      3
f9      3
f10     5
f11     3
f12     3
f13     6
f14     5
f15     0
f16     5
f17     4
f18     2
f19     3
f20     6
target   0
dtype: int64

```

```

[93]: # THE FOLLOWING CODE WAS GENERATED WITH THE HELP OF CHATGPT, THE PROMPTS ARE
      ↵BELOW AND THIS WAS DONE SO THAT I CAN QUICKLY SEE THIS INFORMATION
      # AS I WANTED TO SEE HOW THE DATA BEHAVED AND WHETHER I WOULD NEED TO SCALE THE
      ↵DATA

#prompt:
'''
write me some python code that takes in a pandas dataframe and for each column
    ↵in the dataframe,
makes a graph with the max, min, average, and variance for that column

can you make it so it displays all this info on one graph
'''

#Creates one grouped bar chart displaying max, min, mean, and variance for each
    ↵numeric column in the DataFrame.
def plot_all_column_statistics(df):
    numeric_cols = df.select_dtypes(include='number').columns
    stats = { # Collect statistics
        "Max": [],
        "Min": [],
        "Mean": [],
        "Variance": []}
    for col in numeric_cols:
        series = df[col].dropna()
        stats["Max"].append(series.max())
        stats["Min"].append(series.min())
        stats["Mean"].append(series.mean())
        stats["Variance"].append(series.var())
    # Plotting
    x = np.arange(len(numeric_cols)) # column positions
    width = 0.2 # bar width

```


```

plt.figure(figsize=(12, 6))
plt.bar(x - 1.5*width, stats["Max"], width, label="Max")
plt.bar(x - 0.5*width, stats["Min"], width, label="Min")
plt.bar(x + 0.5*width, stats["Mean"], width, label="Mean")
plt.bar(x + 1.5*width, stats["Variance"], width, label="Variance")

plt.xticks(x, numeric_cols)
plt.ylabel("Value")
plt.title("Column Statistics (Max, Min, Mean, Variance)")
plt.legend()
plt.tight_layout()
plt.show()

```

[94]: df\_testing = df\_train.drop('target', axis=1)  
plot\_all\_column\_statistics(df\_testing)



Notes:

The above graph again shows that the data seems to be mostly on the same scale but with some of these extreme max and mins, scaling the data would be helpful

[95]: X\_train = df\_train.drop('target', axis=1) #spliting the data  
y\_train = df\_train['target']

## 2 DATA PREPROCESSOR

This will be a preprocessor used through all of the models.

As such its not going to do anything fancy, just impute and scale the data, for more fancy stuff like using PCA, thatll be defined at that specifc model.

Both of these preprocessors are needed for reasons that are shown above

```
[96]: cols = ['f1', 'f2', 'f3', 'f4', 'f5', 'f6', 'f7', 'f8', 'f9', 'f10', 'f11', 'f12', 'f13', 'f14', 'f15', 'f16', 'f17', 'f18', 'f19', 'f20']

from sklearn.preprocessing import StandardScaler,PolynomialFeatures
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import FunctionTransformer
from sklearn.compose import make_column_selector as selector

feature_processor=Pipeline(steps = [
    ('imputer',SimpleImputer(strategy='median')),
    ('scaler',StandardScaler())
])

processor = ColumnTransformer(transformers = [
    ('processor',feature_processor,cols)],remainder='passthrough' )
processor
```

```
[96]: ColumnTransformer(remainder='passthrough',
                        transformers=[('processor',
                                       Pipeline(steps=[('imputer',
                                                       SimpleImputer(strategy='median')),
                                                       ('scaler', StandardScaler())])),
                                       ['f1', 'f2', 'f3', 'f4', 'f5', 'f6', 'f7',
                                         'f8', 'f9', 'f10', 'f11', 'f12', 'f13', 'f14',
                                         'f15', 'f16', 'f17', 'f18', 'f19', 'f20'])])
```

### 3 SOME FUNCTION DEFS FOR CONVIENCE

```
[97]: from sklearn.metrics import mean_squared_error
=======
#This funciton splits the data into training and validation sets and plots the
#performance of the mode
=======
def make_train_vs_test_curves(model,X_train,y_train):
    train_size_abs, train_scores, test_scores = learning_curve(
        model, X_train, y_train, train_sizes=[0.3, 0.6, 0.9]
    )
    plt.plot(train_scores.flatten(),color='red')
    plt.plot(test_scores.flatten(),color='blue')
=======

```

```

#This function shows the orginal target data on top of the predicted data based
#on the trained model, while also outputting the square root of the MSE score.
=====

def training_data_fit_display(model,X_train,y_train):
    x_train_pred = model.predict(X_train)
    print("Here is the MSE score:")
    print(np.sqrt(mean_squared_error(y_train,x_train_pred)))
    plt.plot(x_train_pred)
    plt.plot(y_train,color='red')

time_taken_to_train = []
time_taken_to_predict = []

```

## 4 General method of training these models:

Because the training set is, relativly small, I decide to not explicitly seperate out a testing set from the avaiable training data, instead, my general method was to us GridSearchCV to search for hyperparamters and only after finding these params, performing train test splitting to train and then validate the model on the testing set, doing this several times. (this was achieved with the make\_train\_vs\_test\_curves function defined above).

Additionally beacuse I wanted to avoid overfitting, I began without using polynomial features for each model, only adding them after if I decided they had performed well enough to continue onwards.

My process for choosing the best models was based on the Kaggle submission score. I knew that overfitting could be a problem so all models i submitted to kaggle at least once (even if they didnt have the best scores from my metrics, this was because something could have had a good train MSE score, but being overfit it does bad on the testing data, were as something without the best scores could still have done well in the testing data). Then based on those scores i either dropped the model, or tried to continue and refine it. (This can be seen with the ridge and lasso regression models, which performed ok with my metrics, but did the best in kaggle, so i made more versions/iterations of them till they got better)

Finally to avoid confusion I trained models 1-8 in that order, after seeing there performance on kaggle and on my metrics, i then choose Rideg and Lasso to do further iteraions on with models the the .2 versions and etc

## 5 MODEL 1: REGRESSION

No regularization and no polynomial features, so nothing to grid search

```
[98]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

test_pipeline = Pipeline(steps=[
    ('processor',feature_processor),
```


```
('model',LinearRegression())
])

import time

start = time.perf_counter()
test_pipeline.fit(X_train,y_train.to_numpy())
end = time.perf_counter()
time_taken_to_train.append(end-start)

start = time.perf_counter()
test_pipeline.predict(X_train)
end = time.perf_counter()
time_taken_to_predict.append(end-start)
```


```
[99]: make_train_vs_test_curves(test_pipeline,X_train,y_train)
```



```
[100]: training_data_fit_display(test_pipeline,X_train,y_train)
```

Here is the MSE score:

6.213200464581835



*Summary:*

The model performs ok, the MSE score isn't the best, and more importantly the testing score in the validation doesn't perform the best

## 6 MODEL 2: SVM

Just SVM, nothing on top of it

```
[101]: from sklearn.svm import SVR

test_pipeline = Pipeline(steps=[
    ('processor', feature_processor),
    ('model', SVR())
])

params_dict = {
    "model__C": np.linspace(1, 10, 30),
    "model__epsilon": np.linspace(.01, 1, 20) #the two SVM metrics
}
grid = GridSearchCV(test_pipeline,
                    params_dict,
```

```

        cv = 10,
        scoring = 'neg_mean_squared_error',
        verbose = 1,
        n_jobs = -1
    )

grid.fit(X_train,y_train)

```


Fitting 10 folds for each of 600 candidates, totalling 6000 fits

```
[101]: GridSearchCV(cv=10,
                     estimator=Pipeline(steps=[('processor',
                                                Pipeline(steps=[('imputer',
                                                               SimpleImputer(strategy='median')),
                                                               ('scaler',
                                                               StandardScaler()))]),
                                                ('model', SVR())),
                     n_jobs=-1,
                     param_grid={'model__C': array([ 1.           ,  1.31034483,
                     1.62068966,  1.93103448,  2.24137931,
                     2.55172414,  2.86206897,  3.17241379,  3.48275862,  3.79310345,
                     4.10344828,  4.4137931 ,  4.72413793,  5.034482...
                     7.20689655,  7.51724138,  7.82758621,  8.13793103,  8.44827586,
                     8.75862069,  9.06896552,  9.37931034,  9.68965517, 10.          ],
                     'model__epsilon': array([0.01       , 0.06210526,
                     0.11421053, 0.16631579, 0.21842105,
                     0.27052632, 0.32263158, 0.37473684, 0.42684211, 0.47894737,
                     0.53105263, 0.58315789, 0.63526316, 0.68736842, 0.73947368,
                     0.79157895, 0.84368421, 0.89578947, 0.94789474, 1.          ])},
                     scoring='neg_mean_squared_error', verbose=1)
```

```
[102]: grid.best_params_
start = time.perf_counter()
grid.predict(X_train)
end = time.perf_counter()
time_taken_to_predict.append(end-start)
```


```
[103]: training_data_fit_display(grid,X_train,y_train)
```

Here is the MSE score:  
5.252533495200473



```
[104]: best_svm = grid.best_estimator_
make_train_vs_test_curves(best_svm,X_train,y_train)

start = time.perf_counter()
best_svm.fit(X_train,y_train)
end = time.perf_counter()
time_taken_to_train.append(end-start)
```



*Summary:*

This model performs quite well, with a consistently higher testing score over training in validation, along with a good MSE.

However, when submitted to Kaggle the model didn't perform as well as I hoped. I believe that it was either overfitting or a lack of model complexity.

## 7 MODEL 3: RANDOM FOREST REGRESSION

```
[105]: from sklearn.ensemble import RandomForestRegressor
```


```
[106]: test_pipeline = Pipeline(steps=[  
    ('processor', feature_processor),  
    ('model', RandomForestRegressor())  
)
```

```
[107]: from sklearn.model_selection import RandomizedSearchCV  
  
param_dic = {  
    'model__n_estimators': range(20,250),  
    'model__max_depth': range(2,10),
```

```

'model__min_samples_split': range(10,250), #the metrics used for random forest reg
'model__min_samples_leaf': range(5,250)
}
grid = RandomizedSearchCV(test_pipeline,
                           param_dic,
                           n_iter=50,
                           scoring='neg_mean_squared_error',
                           n_jobs=-1,
                           cv=5
)
grid.fit(X_train,y_train)
best_random_forest_reg = grid.best_estimator_
make_train_vs_test_curves(best_random_forest_reg,X_train,y_train)

```



```
[108]: training_data_fit_display(best_random_forest_reg,X_train,y_train)
start = time.perf_counter()
best_random_forest_reg.predict(X_train)
end = time.perf_counter()
time_taken_to_predict.append(end-start)
```


```

start = time.perf_counter()
best_random_forest_reg.fit(X_train,y_train)
end = time.perf_counter()
time_taken_to_train.append(end-start)

```

Here is the MSE score:

6.06620035143386



*Summary:*

This model didn't do quite well, the MSE score was bad, the train test validation returned poor results and most importantly submitting to kaggle returned abysmal results, so I'm going to drop pursuing this model.

## 8 MODEL 4: RIDGE REGRESSION WITH POLYNOMIAL FEATURES

[109]: `from sklearn.linear_model import Ridge`

```

test_pipeline = Pipeline(steps=[
    ('processor',feature_processor),
    ('model',Ridge())
]


```

```

])
params_dict = {
    "model__alpha":np.linspace(0,1,50),
}
grid = GridSearchCV(test_pipeline,
                     params_dict,
                     cv = 10,
                     scoring = 'neg_mean_absolute_error',
                     verbose = 1,
                     n_jobs = -1
)
grid.fit(X_train,y_train)
best_ridge = grid.best_estimator_
make_train_vs_test_curves(best_ridge,X_train,y_train)

```

Fitting 10 folds for each of 50 candidates, totalling 500 fits



```

[110]: training_data_fit_display(best_ridge,X_train,y_train)
start = time.perf_counter()
best_ridge.predict(X_train)
end = time.perf_counter()

```

```


time_taken_to_predict.append(end-start)

start = time.perf_counter()
best_ridge.fit(X_train,y_train)
end = time.perf_counter()
time_taken_to_train.append(end-start)

```

Here is the MSE score:

6.213283783436817



*Summary:*

This model, while initially not seeming to perform as well as SVM in MSE, and train test validation, it seems this model was much better at generalizing, because submitting it to kaggle produced quite good results, as such, a second round of training was performed in an attempt to better nail down the ideal params

## 9 MODEL 4.2 MORE TRAINING

Same setup as above, but more time was taken to train, with a wider range of values. I also manually narrowed the range of params over repeated iterations based on what the grid search would converge to so that I could get even better values

```
[111]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.decomposition import PCA
test_pipeline = Pipeline(steps=[
    ('processor',feature_processor),
    ('poly',PolynomialFeatures()),
    ('model',Ridge())
])

params_dict = {
    "model__alpha":np.linspace(1100,1200,500), #within this range is what is
    ↪eventually settled on as being ideal
    "poly__degree":[3]
}
grid = GridSearchCV(test_pipeline,
                     params_dict,
                     cv = 10,
                     scoring = 'neg_mean_absolute_error',
                     verbose = 1,
                     n_jobs = -1
)


grid.fit(X_train,y_train)
best_ridge_2 = grid.best_estimator_

start = time.perf_counter()
best_ridge_2.fit(X_train,y_train)
end = time.perf_counter()
time_taken_to_train.append(end-start)

start = time.perf_counter()
grid.predict(X_train)
end = time.perf_counter()
time_taken_to_predict.append(end-start)

make_train_vs_test_curves(best_ridge_2,X_train,y_train)
```


Fitting 10 folds for each of 500 candidates, totalling 5000 fits



```
[112]: training_data_fit_display(best_ridge_2,X_train,y_train)
```

Here is the MSE score:

3.151414116911473



```
[113]: grid.best_params_ #these are the params that eventually was settled on
```

```
[113]: {'model__alpha': np.float64(1154.308617234469), 'poly__degree': 3}
```

*Summary:*

The performance here was great, not only were my metrics improved (mainly the train test validation) but the kaggle performance improved as well, so in an attempt to gain better performance, because I seem to have made the params as ideal as I can, I'm going to try using PCA.

## 10 MODEL 4.3 SAME AS ABOVE BUT THIS TIME WITH PCA

```
[114]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.decomposition import PCA
test_pipeline = Pipeline(steps=[
    ('processor', feature_processor),
    ('pca', PCA()),
    ('poly', PolynomialFeatures()),
    ('model', Ridge())
])
```

```

params_dict = {
    "model__alpha":np.linspace(1300,1400,500),
    "poly__degree": [3],
    "pca__n_components": [17] #these are only after several training goes which
    ↪ seemed to have the grid search converging towards these values
}

grid = GridSearchCV(test_pipeline,
                     params_dict,
                     cv = 10,
                     scoring = 'neg_mean_absolute_error',
                     verbose = 1,
                     n_jobs = -1
)
grid.fit(X_train,y_train)
best_ridge_3 = grid.best_estimator_


start = time.perf_counter()
best_ridge_3.fit(X_train,y_train)
end = time.perf_counter()
time_taken_to_train.append(end-start)

start = time.perf_counter()
grid.predict(X_train)
end = time.perf_counter()
time_taken_to_predict.append(end-start)

make_train_vs_test_curves(best_ridge_3,X_train,y_train)

```


Fitting 10 folds for each of 500 candidates, totalling 5000 fits



```
[115]: training_data_fit_display(best_ridge_3,X_train,y_train)
```

Here is the MSE score:

3.472270916747489



```
[116]: grid.best_params_ #best values
```

```
[116]: {'model__alpha': np.float64(1306.8136272545091),
 'pca__n_components': 17,
 'poly__degree': 3}
```

*Summary:*

The performance of this model was NOT improved by PCA, it seems that all the dimensions of the data contribute, because even the reduction of 3 dims hurt quite badly the performance of the model, as such im going to not use PCA at all from here on out.

## 11 MODEL 5 DECISION TREE REGRESSION

```
[117]: from sklearn.tree import DecisionTreeRegressor
```


```
test_pipeline = Pipeline(steps=[  
    ('processor', feature_processor),  
    ('model', DecisionTreeRegressor())  
)  
  
params_dict={
```

```

'model__max_depth':range(1,5),
'model__min_samples_split':range(5,50,2),
'model__min_samples_leaf':range(5,50,2),
'model__max_leaf_nodes':[None,3,5,10],
}
grid = GridSearchCV(test_pipeline,
                     params_dict,
                     cv = 10,
                     scoring = 'neg_mean_absolute_error',
                     verbose = 1,
                     n_jobs = -1
)
grid.fit(X_train,y_train)
best_tree = grid.best_estimator_
make_train_vs_test_curves(best_ridge,X_train,y_train)

```

Fitting 10 folds for each of 8464 candidates, totalling 84640 fits



```
[118]: training_data_fit_display(best_tree,X_train,y_train)
start = time.perf_counter()
best_tree.predict(X_train)
end = time.perf_counter()
```

```


time_taken_to_predict.append(end-start)

start = time.perf_counter()
best_tree.fit(X_train,y_train)
end = time.perf_counter()
time_taken_to_train.append(end-start)

```

Here is the MSE score:

7.4110216300609



*Summary:*

This model performed at the end of the day quite bad, compared to other models. This makes sense though, if a random forest performed poorly, then a regular decision tree is also more likely to perform poorly as well.

## 12 MODEL 6 LASSO

Without polynomial features

```
[119]: from sklearn.linear_model import Lasso
```

```
test_pipeline = Pipeline(steps=[
```


```

        ('processor',feature_processor),
        ('model',Lasso())
    ])

params_dict = {
    "model__alpha":np.linspace(0.001,1,75),
    "model__fit_intercept":[False,True]
}
grid = GridSearchCV(test_pipeline,
                     params_dict,
                     cv = 10,
                     scoring = 'neg_mean_absolute_error',
                     verbose = 1,
                     n_jobs = -1
)
grid.fit(X_train,y_train)
best_lasso = grid.best_estimator_
make_train_vs_test_curves(best_lasso,X_train,y_train)

```

Fitting 10 folds for each of 150 candidates, totalling 1500 fits




```
[120]: training_data_fit_display(best_lasso,X_train,y_train)
start = time.perf_counter()
best_lasso.predict(X_train)
end = time.perf_counter()
time_taken_to_predict.append(end-start)

start = time.perf_counter()
best_lasso.fit(X_train,y_train)
end = time.perf_counter()
time_taken_to_train.append(end-start)
```

Here is the MSE score:

6.330678391528587



*Summary:*

This model performed surprisingly well. While the metrics I tested it on (seen above) didn't show AMAZING performance, the kaggle score was the best that I had gotten up to this point, so I continued to follow done this model type.

## 13 MODEL 6.2 MORE TRAINING ON THE LASSO

This time with polynomial features

As I did on Model 4.2 (ridge reg with polynomial features) i manuelli narrowed the range of params that i used based on what the model was converging towards to get better results

```
[121]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.decomposition import PCA

test_pipeline = Pipeline(steps=[
    ('processor',feature_processor),
    ('poly',PolynomialFeatures()),
    ('model',Lasso())
])

params_dict = {
    "model__alpha":np.linspace(.25,.75,100),
    "poly__degree":[3,4],
    "model__max_iter":[5000]
}
grid = GridSearchCV(test_pipeline,
                     params_dict,
                     cv = 10,
                     scoring = 'neg_mean_absolute_error',
                     verbose = 1,
                     n_jobs = -1
                    )


grid.fit(X_train,y_train)
best_lasso_2 = grid.best_estimator_

start = time.perf_counter()
best_lasso_2.fit(X_train,y_train)
end = time.perf_counter()
time_taken_to_train.append(end-start)

start = time.perf_counter()
grid.predict(X_train)
end = time.perf_counter()
time_taken_to_predict.append(end-start)

make_train_vs_test_curves(best_lasso_2,X_train,y_train)
```


Fitting 10 folds for each of 200 candidates, totalling 2000 fits



```
[122]: training_data_fit_display(best_lasso_2,X_train,y_train)
```

Here is the MSE score:

4.044953021119381



```
[123]: grid.best_params_ #final best params
```

```
[123]: {'model__alpha': np.float64(0.5025252525252526),
 'model__max_iter': 5000,
 'poly__degree': 3}
```

*Summary:*

This was the best model that I could find, performing miles better then any other model in Kaggle along with the train test validation scores, as such this is my final best Model.

## 14 MODEL 7 ELASTIC NET

Without polynomial features

```
[124]: from sklearn.linear_model import ElasticNet
test_pipeline = Pipeline(steps=[
    ('processor', feature_processor),
    ('model', ElasticNet())
])

params_dict = {
    "model__alpha": np.linspace(.02, 1, 75),
```

```

"model__l1_ratio":np.linspace(.02,1,75),
"model__fit_intercept":[False,True]
}

grid = GridSearchCV(test_pipeline,
                     params_dict,
                     cv = 10,
                     scoring = 'neg_mean_absolute_error',
                     verbose = 1,
                     n_jobs = -1
                    )

grid.fit(X_train,y_train)
best_elastic = grid.best_estimator_


start = time.perf_counter()
best_elastic.fit(X_train,y_train)
end = time.perf_counter()
time_taken_to_train.append(end-start)

start = time.perf_counter()
grid.predict(X_train)
end = time.perf_counter()
time_taken_to_predict.append(end-start)

make_train_vs_test_curves(best_elastic,X_train,y_train)

```


Fitting 10 folds for each of 11250 candidates, totalling 112500 fits



```
[125]: training_data_fit_display(best_elastic,X_train,y_train)
```

Here is the MSE score:

6.411331395949554



*Summary:*

After the performance of the ridge and lasso i figured testing both combined would be ideal, however it didnt pruduce the best results, being out performed but both ridge and lasso regresssion individually, so i abandoned this route for following ridge and lasso by themselves

## 15 MODEL 8 NN

This one i manuelli narrowed the range of how many layers to use based on kaggle performance. I didnt have high hopes for this model type, and it didnt perform saddly.

```
[126]: from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout
from numpy.random import seed
import tensorflow
tensorflow.random.set_seed(42)

y_train_np = feature_processor.fit_transform(pd.DataFrame(y_train))
X_train_np = feature_processor.fit_transform(X_train)
```

```
[127]: from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential([
    layers.Dense(10, activation='relu', input_shape=(20,)),
    layers.Dense(1) # Output layer for regression
])
model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])
history = model.fit(X_train_np, y_train_np, epochs=100, batch_size=32, validation_split=0.2)
```

Epoch 1/100

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/keras/src/layers/core/dense.py:92: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

10/10 0s 10ms/step -  
loss: 1.4544 - mae: 0.9397 - val\_loss: 1.5314 - val\_mae: 0.9449

Epoch 2/100

10/10 0s 3ms/step - loss:  
1.3135 - mae: 0.8829 - val\_loss: 1.3798 - val\_mae: 0.8986

Epoch 3/100

10/10 0s 3ms/step - loss:  
1.1979 - mae: 0.8329 - val\_loss: 1.2512 - val\_mae: 0.8558

Epoch 4/100

10/10 0s 3ms/step - loss:  
1.1007 - mae: 0.7901 - val\_loss: 1.1422 - val\_mae: 0.8169

Epoch 5/100

10/10 0s 3ms/step - loss:  
1.0196 - mae: 0.7532 - val\_loss: 1.0501 - val\_mae: 0.7826

Epoch 6/100

10/10 0s 3ms/step - loss:  
0.9517 - mae: 0.7215 - val\_loss: 0.9724 - val\_mae: 0.7522

Epoch 7/100

10/10 0s 7ms/step - loss:  
0.8944 - mae: 0.6944 - val\_loss: 0.9069 - val\_mae: 0.7245

Epoch 8/100

10/10 0s 3ms/step - loss:  
0.8451 - mae: 0.6703 - val\_loss: 0.8500 - val\_mae: 0.6989

Epoch 9/100

10/10 0s 3ms/step - loss:  
0.8022 - mae: 0.6486 - val\_loss: 0.7990 - val\_mae: 0.6753

Epoch 10/100

10/10 0s 3ms/step - loss:  
0.7643 - mae: 0.6285 - val\_loss: 0.7526 - val\_mae: 0.6536

Epoch 11/100

```
10/10          0s 3ms/step - loss:  
0.7305 - mae: 0.6099 - val_loss: 0.7130 - val_mae: 0.6334  
Epoch 12/100  
10/10          0s 3ms/step - loss:  
0.6991 - mae: 0.5927 - val_loss: 0.6778 - val_mae: 0.6153  
Epoch 13/100  
10/10          0s 3ms/step - loss:  
0.6711 - mae: 0.5775 - val_loss: 0.6466 - val_mae: 0.5985  
Epoch 14/100  
10/10          0s 3ms/step - loss:  
0.6460 - mae: 0.5642 - val_loss: 0.6186 - val_mae: 0.5832  
Epoch 15/100  
10/10          0s 3ms/step - loss:  
0.6235 - mae: 0.5524 - val_loss: 0.5936 - val_mae: 0.5693  
Epoch 16/100  
10/10          0s 3ms/step - loss:  
0.6030 - mae: 0.5418 - val_loss: 0.5710 - val_mae: 0.5576  
Epoch 17/100  
10/10          0s 3ms/step - loss:  
0.5844 - mae: 0.5318 - val_loss: 0.5501 - val_mae: 0.5477  
Epoch 18/100  
10/10          0s 3ms/step - loss:  
0.5673 - mae: 0.5228 - val_loss: 0.5310 - val_mae: 0.5381  
Epoch 19/100  
10/10          0s 3ms/step - loss:  
0.5512 - mae: 0.5144 - val_loss: 0.5135 - val_mae: 0.5290  
Epoch 20/100  
10/10          0s 3ms/step - loss:  
0.5364 - mae: 0.5064 - val_loss: 0.4975 - val_mae: 0.5210  
Epoch 21/100  
10/10          0s 3ms/step - loss:  
0.5226 - mae: 0.4989 - val_loss: 0.4831 - val_mae: 0.5133  
Epoch 22/100  
10/10          0s 3ms/step - loss:  
0.5099 - mae: 0.4917 - val_loss: 0.4698 - val_mae: 0.5063  
Epoch 23/100  
10/10          0s 3ms/step - loss:  
0.4979 - mae: 0.4848 - val_loss: 0.4577 - val_mae: 0.5002  
Epoch 24/100  
10/10          0s 3ms/step - loss:  
0.4866 - mae: 0.4784 - val_loss: 0.4466 - val_mae: 0.4945  
Epoch 25/100  
10/10          0s 3ms/step - loss:  
0.4760 - mae: 0.4724 - val_loss: 0.4366 - val_mae: 0.4889  
Epoch 26/100  
10/10          0s 3ms/step - loss:  
0.4661 - mae: 0.4670 - val_loss: 0.4274 - val_mae: 0.4840  
Epoch 27/100
```

```
10/10          0s 3ms/step - loss:  
0.4566 - mae: 0.4620 - val_loss: 0.4189 - val_mae: 0.4791  
Epoch 28/100  
10/10          0s 3ms/step - loss:  
0.4477 - mae: 0.4573 - val_loss: 0.4109 - val_mae: 0.4742  
Epoch 29/100  
10/10          0s 3ms/step - loss:  
0.4393 - mae: 0.4530 - val_loss: 0.4036 - val_mae: 0.4694  
Epoch 30/100  
10/10          0s 3ms/step - loss:  
0.4313 - mae: 0.4487 - val_loss: 0.3968 - val_mae: 0.4649  
Epoch 31/100  
10/10          0s 3ms/step - loss:  
0.4236 - mae: 0.4446 - val_loss: 0.3905 - val_mae: 0.4610  
Epoch 32/100  
10/10          0s 3ms/step - loss:  
0.4164 - mae: 0.4406 - val_loss: 0.3843 - val_mae: 0.4570  
Epoch 33/100  
10/10          0s 3ms/step - loss:  
0.4094 - mae: 0.4366 - val_loss: 0.3786 - val_mae: 0.4532  
Epoch 34/100  
10/10          0s 3ms/step - loss:  
0.4027 - mae: 0.4327 - val_loss: 0.3733 - val_mae: 0.4498  
Epoch 35/100  
10/10          0s 3ms/step - loss:  
0.3962 - mae: 0.4290 - val_loss: 0.3682 - val_mae: 0.4464  
Epoch 36/100  
10/10          0s 3ms/step - loss:  
0.3899 - mae: 0.4257 - val_loss: 0.3637 - val_mae: 0.4433  
Epoch 37/100  
10/10          0s 3ms/step - loss:  
0.3839 - mae: 0.4225 - val_loss: 0.3594 - val_mae: 0.4402  
Epoch 38/100  
10/10          0s 3ms/step - loss:  
0.3781 - mae: 0.4193 - val_loss: 0.3556 - val_mae: 0.4372  
Epoch 39/100  
10/10          0s 3ms/step - loss:  
0.3726 - mae: 0.4164 - val_loss: 0.3521 - val_mae: 0.4345  
Epoch 40/100  
10/10          0s 3ms/step - loss:  
0.3674 - mae: 0.4136 - val_loss: 0.3490 - val_mae: 0.4320  
Epoch 41/100  
10/10          0s 3ms/step - loss:  
0.3624 - mae: 0.4109 - val_loss: 0.3462 - val_mae: 0.4295  
Epoch 42/100  
10/10          0s 3ms/step - loss:  
0.3577 - mae: 0.4085 - val_loss: 0.3436 - val_mae: 0.4272  
Epoch 43/100
```

```
10/10          0s 3ms/step - loss:  
0.3531 - mae: 0.4063 - val_loss: 0.3413 - val_mae: 0.4250  
Epoch 44/100  
10/10          0s 3ms/step - loss:  
0.3488 - mae: 0.4041 - val_loss: 0.3390 - val_mae: 0.4229  
Epoch 45/100  
10/10          0s 3ms/step - loss:  
0.3446 - mae: 0.4019 - val_loss: 0.3372 - val_mae: 0.4209  
Epoch 46/100  
10/10          0s 3ms/step - loss:  
0.3404 - mae: 0.3997 - val_loss: 0.3356 - val_mae: 0.4191  
Epoch 47/100  
10/10          0s 3ms/step - loss:  
0.3364 - mae: 0.3976 - val_loss: 0.3340 - val_mae: 0.4175  
Epoch 48/100  
10/10          0s 3ms/step - loss:  
0.3326 - mae: 0.3956 - val_loss: 0.3324 - val_mae: 0.4162  
Epoch 49/100  
10/10          0s 3ms/step - loss:  
0.3288 - mae: 0.3937 - val_loss: 0.3308 - val_mae: 0.4149  
Epoch 50/100  
10/10          0s 3ms/step - loss:  
0.3253 - mae: 0.3918 - val_loss: 0.3291 - val_mae: 0.4134  
Epoch 51/100  
10/10          0s 3ms/step - loss:  
0.3218 - mae: 0.3900 - val_loss: 0.3276 - val_mae: 0.4119  
Epoch 52/100  
10/10          0s 3ms/step - loss:  
0.3185 - mae: 0.3883 - val_loss: 0.3263 - val_mae: 0.4110  
Epoch 53/100  
10/10          0s 3ms/step - loss:  
0.3154 - mae: 0.3867 - val_loss: 0.3249 - val_mae: 0.4100  
Epoch 54/100  
10/10          0s 3ms/step - loss:  
0.3124 - mae: 0.3851 - val_loss: 0.3237 - val_mae: 0.4090  
Epoch 55/100  
10/10          0s 3ms/step - loss:  
0.3095 - mae: 0.3835 - val_loss: 0.3224 - val_mae: 0.4083  
Epoch 56/100  
10/10          0s 3ms/step - loss:  
0.3068 - mae: 0.3820 - val_loss: 0.3215 - val_mae: 0.4077  
Epoch 57/100  
10/10          0s 3ms/step - loss:  
0.3042 - mae: 0.3806 - val_loss: 0.3205 - val_mae: 0.4070  
Epoch 58/100  
10/10          0s 3ms/step - loss:  
0.3017 - mae: 0.3790 - val_loss: 0.3195 - val_mae: 0.4062  
Epoch 59/100
```

```
10/10          0s 3ms/step - loss:  
0.2992 - mae: 0.3775 - val_loss: 0.3184 - val_mae: 0.4055  
Epoch 60/100  
10/10          0s 3ms/step - loss:  
0.2969 - mae: 0.3760 - val_loss: 0.3175 - val_mae: 0.4049  
Epoch 61/100  
10/10          0s 3ms/step - loss:  
0.2946 - mae: 0.3745 - val_loss: 0.3165 - val_mae: 0.4045  
Epoch 62/100  
10/10          0s 3ms/step - loss:  
0.2924 - mae: 0.3730 - val_loss: 0.3158 - val_mae: 0.4044  
Epoch 63/100  
10/10          0s 3ms/step - loss:  
0.2902 - mae: 0.3717 - val_loss: 0.3151 - val_mae: 0.4042  
Epoch 64/100  
10/10          0s 3ms/step - loss:  
0.2882 - mae: 0.3704 - val_loss: 0.3145 - val_mae: 0.4041  
Epoch 65/100  
10/10          0s 3ms/step - loss:  
0.2861 - mae: 0.3691 - val_loss: 0.3139 - val_mae: 0.4039  
Epoch 66/100  
10/10          0s 3ms/step - loss:  
0.2842 - mae: 0.3678 - val_loss: 0.3137 - val_mae: 0.4039  
Epoch 67/100  
10/10          0s 3ms/step - loss:  
0.2823 - mae: 0.3666 - val_loss: 0.3136 - val_mae: 0.4042  
Epoch 68/100  
10/10          0s 3ms/step - loss:  
0.2804 - mae: 0.3654 - val_loss: 0.3134 - val_mae: 0.4043  
Epoch 69/100  
10/10          0s 3ms/step - loss:  
0.2786 - mae: 0.3640 - val_loss: 0.3130 - val_mae: 0.4043  
Epoch 70/100  
10/10          0s 3ms/step - loss:  
0.2769 - mae: 0.3627 - val_loss: 0.3127 - val_mae: 0.4044  
Epoch 71/100  
10/10          0s 3ms/step - loss:  
0.2753 - mae: 0.3615 - val_loss: 0.3124 - val_mae: 0.4044  
Epoch 72/100  
10/10          0s 3ms/step - loss:  
0.2736 - mae: 0.3604 - val_loss: 0.3119 - val_mae: 0.4042  
Epoch 73/100  
10/10          0s 3ms/step - loss:  
0.2721 - mae: 0.3591 - val_loss: 0.3114 - val_mae: 0.4040  
Epoch 74/100  
10/10          0s 3ms/step - loss:  
0.2706 - mae: 0.3581 - val_loss: 0.3111 - val_mae: 0.4040  
Epoch 75/100
```

```
10/10          0s 3ms/step - loss:  
0.2692 - mae: 0.3571 - val_loss: 0.3107 - val_mae: 0.4038  
Epoch 76/100  
10/10          0s 3ms/step - loss:  
0.2677 - mae: 0.3560 - val_loss: 0.3104 - val_mae: 0.4037  
Epoch 77/100  
10/10          0s 3ms/step - loss:  
0.2663 - mae: 0.3549 - val_loss: 0.3099 - val_mae: 0.4033  
Epoch 78/100  
10/10          0s 3ms/step - loss:  
0.2649 - mae: 0.3538 - val_loss: 0.3097 - val_mae: 0.4033  
Epoch 79/100  
10/10          0s 3ms/step - loss:  
0.2636 - mae: 0.3529 - val_loss: 0.3096 - val_mae: 0.4034  
Epoch 80/100  
10/10          0s 6ms/step - loss:  
0.2622 - mae: 0.3519 - val_loss: 0.3093 - val_mae: 0.4033  
Epoch 81/100  
10/10          0s 3ms/step - loss:  
0.2609 - mae: 0.3509 - val_loss: 0.3092 - val_mae: 0.4032  
Epoch 82/100  
10/10          0s 3ms/step - loss:  
0.2597 - mae: 0.3499 - val_loss: 0.3089 - val_mae: 0.4031  
Epoch 83/100  
10/10          0s 3ms/step - loss:  
0.2585 - mae: 0.3489 - val_loss: 0.3086 - val_mae: 0.4029  
Epoch 84/100  
10/10          0s 3ms/step - loss:  
0.2573 - mae: 0.3480 - val_loss: 0.3087 - val_mae: 0.4029  
Epoch 85/100  
10/10          0s 3ms/step - loss:  
0.2562 - mae: 0.3471 - val_loss: 0.3085 - val_mae: 0.4027  
Epoch 86/100  
10/10          0s 3ms/step - loss:  
0.2551 - mae: 0.3461 - val_loss: 0.3084 - val_mae: 0.4025  
Epoch 87/100  
10/10          0s 3ms/step - loss:  
0.2540 - mae: 0.3453 - val_loss: 0.3082 - val_mae: 0.4023  
Epoch 88/100  
10/10          0s 3ms/step - loss:  
0.2529 - mae: 0.3444 - val_loss: 0.3081 - val_mae: 0.4022  
Epoch 89/100  
10/10          0s 3ms/step - loss:  
0.2518 - mae: 0.3436 - val_loss: 0.3083 - val_mae: 0.4022  
Epoch 90/100  
10/10          0s 3ms/step - loss:  
0.2507 - mae: 0.3428 - val_loss: 0.3081 - val_mae: 0.4021  
Epoch 91/100
```

```

10/10          0s 3ms/step - loss:
0.2497 - mae: 0.3419 - val_loss: 0.3082 - val_mae: 0.4022
Epoch 92/100
10/10          0s 3ms/step - loss:
0.2487 - mae: 0.3411 - val_loss: 0.3086 - val_mae: 0.4025
Epoch 93/100
10/10          0s 3ms/step - loss:
0.2477 - mae: 0.3404 - val_loss: 0.3087 - val_mae: 0.4027
Epoch 94/100
10/10          0s 3ms/step - loss:
0.2467 - mae: 0.3397 - val_loss: 0.3087 - val_mae: 0.4027
Epoch 95/100
10/10          0s 3ms/step - loss:
0.2458 - mae: 0.3391 - val_loss: 0.3090 - val_mae: 0.4030
Epoch 96/100
10/10          0s 3ms/step - loss:
0.2449 - mae: 0.3385 - val_loss: 0.3092 - val_mae: 0.4032
Epoch 97/100
10/10          0s 3ms/step - loss:
0.2440 - mae: 0.3378 - val_loss: 0.3093 - val_mae: 0.4034
Epoch 98/100
10/10          0s 3ms/step - loss:
0.2431 - mae: 0.3372 - val_loss: 0.3096 - val_mae: 0.4037
Epoch 99/100
10/10          0s 3ms/step - loss:
0.2422 - mae: 0.3367 - val_loss: 0.3098 - val_mae: 0.4040
Epoch 100/100
10/10          0s 3ms/step - loss:
0.2414 - mae: 0.3361 - val_loss: 0.3099 - val_mae: 0.4041

```

[128]: model.summary()

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 10)	210
dense_7 (Dense)	(None, 1)	11

Total params: 665 (2.60 KB)

Trainable params: 221 (884.00 B)

```
Non-trainable params: 0 (0.00 B)
```

```
Optimizer params: 444 (1.74 KB)
```

```
[129]: feature_processor=Pipeline(steps = [
    ('imputer',SimpleImputer(strategy='median')), #MAY CHANGE THIS TO MEAN I
    ↪THINK
    ('scaler',StandardScaler())
])
df_testing_np =feature_processor.fit_transform(df_test)
df_testing_np.shape
```

```
[129]: (800, 20)
```

```
[130]: nn_pred = model.predict(df_testing_np)
```

```
25/25          0s 482us/step
```

*Summary:*


Submission to kaggle showed that this model performed the WORST. Quite disappointing but oh well, and that was after manuelly decreasing the number of layers it still performed the worst in kaggle. Oh well, im going to drop this model then.

Also i know i didnt display any of the model performance metrics, but thats because i didnt want to adapt my code to tensorflow code, and because it performed so bad in kaggle that its not even worth it.

## 16 QUICK INTERLUDE: PREDICTION AND TRAINING TIME


```
[131]: model_names = ['Linear','SVM','Random'
    ↪Forest, 'Ridge', 'Ridge_2', 'Ridge_3', 'Decision'
    ↪Tree, 'Lasso', 'Lasso_2', 'ElasticNet']
plt.bar(model_names,time_taken_to_predict)
plt.xticks(rotation='vertical')
plt.xlabel('Models')
plt.ylabel('Time to make a Prediction')
plt.title('Prediction time vs Model')
```

```
[131]: Text(0.5, 1.0, 'Prediction time vs Model')
```



```
[132]: model_names = ['Linear', 'SVM', 'RandomForest',  
                   'Ridge', 'Ridge_2', 'Ridge_3', 'DecisionTree',  
                   'Lasso', 'Lasso_2', 'ElasticNet']  
plt.bar(model_names, time_taken_to_train)  
plt.xticks(rotation='vertical')  
plt.xlabel('Models')  
plt.ylabel('Time to Train')  
plt.title('Train time vs Model')
```

```
[132]: Text(0.5, 1.0, 'Train time vs Model')
```



*Notes:*

Note that what I timed per model was explicitly the training time, NOT the time taken to perform grid search.

Additional It seems that in general the more complex models take both longer to train and longer to make predicts for.

## 17 MODEL SUBMISSIONS

```
[133]: =====
# This function takes in the mprediction and what the file should be called and
# outputs it to a txt for manuelli submission to kaggle
=====
def make_submission_txt(pred,filename):
    index_arr = np.arange(401, 401+pred.size).reshape(-1,1)
```

```

index_arr = index_arr.astype(int).flatten() #makes index array
final_array = np.stack((index_arr,pred),axis=1) #stacks them
df = pd.DataFrame(final_array,columns=['id','target']) #adds name
df['id'] = df['id'].astype(int) #outputs
df.to_csv(filename+'.csv', index=False)
print("The file has been created")

```

## SVM

[134]:

```

best_svm.fit(X_train,y_train)
best_svm_pred = best_svm.predict(df_test)
make_submission_txt(best_svm_pred,"svm1")

```

The file has been created

best\_random\_forest\_reg

[135]:

```

best_random_forest_reg.fit(X_train,y_train)
best_random_forest_reg_pred = best_random_forest_reg.predict(df_test)
make_submission_txt(best_random_forest_reg_pred,"rfr1")

```

The file has been created

best\_ridge

[136]:

```

best_ridge.fit(X_train,y_train)
best_ridge_pred = best_ridge.predict(df_test)
make_submission_txt(best_ridge_pred,"r1")
#NOTES: so far this one has performed the best when submitted

```

The file has been created

best\_tree

[137]:

```

best_tree.fit(X_train,y_train)
best_tree_pred = best_tree.predict(df_test)
make_submission_txt(best_tree_pred,"t1")

```

The file has been created

best\_lasso

[138]:

```

best_lasso.fit(X_train,y_train)
best_lasso_pred = best_lasso.predict(df_test)
make_submission_txt(best_lasso_pred,"lasso1")

```

The file has been created

best\_elastic

[139]:

```

best_elastic.fit(X_train,y_train)
best_elastic_pred = best_elastic.predict(df_test)
make_submission_txt(best_elastic_pred,"elastic1")

```

The file has been created

**best\_ridge\_2**

```
[140]: best_ridge_2.fit(X_train,y_train)
best_ridge_2_pred = best_ridge_2.predict(df_test)
make_submission_txt(best_ridge_2_pred,"r11")
```

The file has been created

**best\_ridge\_3**

```
[141]: best_ridge_3.fit(X_train,y_train)
best_ridge_3_pred = best_ridge_3.predict(df_test)
make_submission_txt(best_ridge_3_pred,"r12")
```

The file has been created

**best\_lasso\_2**

```
[142]: best_lasso_2.fit(X_train,y_train)
best_lasso_2_pred = best_lasso_2.predict(df_test)
make_submission_txt(best_lasso_2_pred,"l12")
```

The file has been created

NN

```
[143]: make_submission_txt(nn_pred.flatten(),"nn5")
```

The file has been created

## 18 THE BEST MODELS:


- A) Lasso with polynomial features: 18.58
- B) Ridge with polynomial features: 32.55
- C) Ridge with polynomial features and PCA: 47.92
- D) Ridge (By itself): 55.22
- E) Lasso (By itself): 56.93

Ranked by performance in the Kaggle submissions with the MSE score displayed (gotten from kaggles system)

Models A and B and C were able to generalize much better then the other 3, with model A, being able to generalize the best, miles ahead of model B

And finally im going to display again how the best model performed on the metrics i tested on.


```
[144]: best_lasso_2.fit(X_train,y_train)
make_train_vs_test_curves(best_lasso_2,X_train,y_train)
```



```
[145]: training_data_fit_display(best_lasso_2,X_train,y_train)
```

Here is the MSE score:

4.044953021119381



Finally this model got a kaggle score of: 18.58684 Which at least at the time of me making/submitting this assgiment had me in 1st place on the leader board for the class, which i think i kind of cool!